

LAUNCH BRIEF • FIRST-PARTY SKILL HARDENING + MCP SERVER

Your team's own skills are dependencies too.

Agent Skills already hardens 1,000+ community skills. This launch brings the same pipeline to the skills your team writes, and adds an MCP server so every agent pulls the current hardened version on its own.

350,000+

skills on public registries within two months

36%

of scanned public skills had at least one security flaw (Snyk ToxicSkills)

72%

of enterprises are using or testing AI agents

Take first-party hardening and MCP from closed beta to GA.

What ships

- **Hardening-as-a-service for first-party skills.**
Teams submit the skills they write. Chainguard reviews, fixes, and returns each one with a HARDENING.md audit log.
- **Custom rulesets.** Add org-specific rules on top of Chainguard's baseline.
- **Agent Skills MCP server.** Agents query and pull the right skill on demand. When a skill is re-hardened, every agent gets the update.

```
skills.cgr.dev/<org>/<skill_name>:<version tag>
```

Launch tier: 2

- Expands a live product to a new artifact class
- Opens the paid enterprise tier of Agent Skills
- Rides the Gartner MQ Leader news for analyst and field air cover

02 • WHO IT'S FOR

AUDIENCE	ROLE IN THE DEAL	WHAT THEY NEED TO HEAR
Platform engineering	Owns the internal skill library and developer tooling	One registry and one install path for every skill the org uses, community and first-party.
AppSec / security leads	Approves the budget and signs off on risk	An audit trail for every change to every skill, and rules that match your risk posture.
Developers	Daily users in Claude Code, Cursor, Copilot, Gemini CLI	Nothing changes in your workflow. The agent pulls the safe version.

03 • THE REAL ALTERNATIVE

TODAY

Skills live everywhere

Slack threads, shared drives, and personal instances. No versioning, access control, or audit trail.

TODAY

Scanners flag and stop

A report lands with a developer who has no time to rewrite the skill.

TODAY

"Internal means safe"

First-party skills still pull scripts, downloads, and broad permissions.

For platform and security teams whose developers write their own agent skills, **Chainguard Agent Skills** is the registry and hardening pipeline that **fixes risky skills before an agent runs them**, and keeps every agent on the current safe version.

05 • MESSAGE PILLARS, EACH WITH PUBLIC PROOF

It fixes the skill

Every hardened skill ships with a HARDENING.md: each rule checked, what was found, what changed, and proof nothing material broke.

[Chainguard blog, June 17, 2026](#)

It stays current

Skills re-harden when the source changes or a rule is added. Through MCP, agents get the new version with no local copy to update.

[Chainguard blog, June 17, 2026](#)

Version discipline for agent behavior

Pin an agent to a skill SHA. Roll back when something breaks. Registry entitlements stay inside your org.

[Chainguard blog, June 17, 2026](#)

Your rules, on top of ours

Custom rulesets extend the baseline, which already targets over-permissioning, obfuscated commands, and credential harvesting.

[Chainguard blog, June 17, 2026](#)

06 • FIELD ENABLEMENT: OBJECTIONS

THEY SAY	WE SAY
"We already scan our skills."	Keep scanning. Chainguard changes what happens next: the skill comes back fixed, with a diff your auditors can read.
"Our skills are internal."	So are the permissions they grant. A first-party skill carries the same supply chain risk as a community one.
"Another tool to roll out?"	Push and pull with chainctl. Developers install a hardened skill the way they install any SKILL.md today.

Proof first, then noise.

T-30 DAYS

- Two closed-beta customers on the record, one platform and one security voice
- Seller one-pager and objection guide
- Analyst briefings, framed on the new Gartner MQ category

LAUNCH DAY

- Blog, docs on Chainguard Academy, pricing page update
- Demo video: harden a first-party skill end to end and read its HARDENING.md
- Email to Containers and Libraries customers

T+30 DAYS

- Customer story with a before-and-after skill inventory
- Webinar with a beta customer
- Read the numbers and reset targets

08 • HOW I'D MEASURE IT

NORTH STAR

Orgs with first-party skills in their private registry

It shows the product became where a team's own skills live.

Inputs

- Time from signup to first hardened first-party skill
- Agents connected through the MCP server
- Beta-to-paid conversion
- Expansion pipeline from existing Containers and Libraries accounts

Every target is a hypothesis until the first month sets a baseline.

SOURCES

Chainguard, "Introducing Chainguard Agent Skills: Because your AI agent shouldn't trust strangers" (Mar 17, 2026). Chainguard, "Chainguard Agent Skills is now open to everyone, with a private registry to manage your internal skills" (Jun 17, 2026): 1,000+ hardened skills, 350,000+ public skills, 36% (Snyk ToxicSkills), 72% of enterprises, HARDENING.md, private registry, closed beta and MCP details. Chainguard, "Everything we announced during AI Readiness Innovation Week" (Jun 25, 2026): Gartner Magic Quadrant for Software Supply Chain Security. Chainguard Academy, Agent Skills documentation.

About this sample. This is a launch brief for a hypothetical GA of features Chainguard announced in closed beta. I wrote it to show how I'd run a launch: one audience per role in the deal, the real alternative, a positioning line, pillars with public proof, and a plan that lands customer proof before launch day. I write a SKILL.md file for every repeatable workflow in my own work and run most of my job through agents and MCP. The layout follows Chainguard's public design tokens. I used public information only. I work at AWS, and this uses no AWS internal information.