


How I'd take the DataNow Developer Platform to market

Data and ML engineers can adopt DataNow with one line of Python.
This is the plan to reach them first.

```
$ pip install "datanow[fsspec]"
```

```
datanow://team/workspace
```



Reach one audience, make the first call fast, and prove it with one customer in 90 days.

DataNow shipped a Python SDK, a Management API, Connect, and an MCP server. Enterprise sales stays the growth engine. This plan adds a self-serve path beside it.

-
- 01 Start with data and ML engineers.**
They already use pandas, Dask, and PyArrow. The SDK's fsspec support means their code changes by one file path.
 - 02 Build a self-serve path next to sales.**
Developers learn from docs and communities. Vendor events reach 3% of them.
 - 03 Cut time to first call from about 10 minutes to under 5.**
A self-serve token and a code-first quickstart close most of the gap.
 - 04 Run on one north star.**
Weekly active accounts running SDK file operations. Phase gates decide what ships next.
-

STARTING POINT

The platform starts inside an installed base that already trusts the product.

3,000+

companies already run DataNow. Every one is a place the platform can grow.

80,000+
users


100+
countries

60+ PB
stored

Sales-led
enterprise motion today

Developers learn from communities and docs. Vendor events reach 3% of them.

Share of developers using each resource to stay up to date



52%

of developers with 10+ years of experience use vendor docs and sites, compared with 32% of early-career developers.

Senior engineers pick infrastructure. They read the docs. So the docs and the code samples do the marketing.

“

*My training data lives on a shared workspace.
I copy it to local disk before every run, then I
wait on the copy.*

Composite ML engineer. A hypothesis to test in the first ten customer interviews.

Copy before every run

Two copies of every dataset

Scripts that break when paths move

Existing pandas code reads from a workspace when one path changes.

The SDK implements fsspec, the storage interface that pandas, Dask, and PyArrow already use. The engineer keeps the code, the libraries, and the habits.

The message writes itself: **your code doesn't change.**


```
train.py

import os
import pandas as pd

# Before: copy the shard to local disk, then read it
df = pd.read_parquet("/mnt/local/train/shard-001.parquet")

# After: read it in place from the workspace
df = pd.read_parquet(
    "datanow://team/workspace/train/shard-001.parquet",
    storage_options={"token": os.environ["DATANOW_TOKEN"]},
)
```

Four entry points share one streaming engine.



Python SDK BETA

The developer's front door. Clientless, so no desktop app is needed.

Management API GA

Automates admin work that happens in the web client today.

Connect GA

Brings existing S3 data into a workspace without a migration.

MCP Server BETA

Lets AI agents find and describe files in plain language.

I made the first call myself. It took about 10 minutes.

Hand-timed on a comparable public SDK with a trial account. I wrote a file and read it back through code.

```
$ pip install datanow
Successfully installed datanow-0.9.0
$ python first_call.py
wrote /hello.txt
read /hello.txt → b'Hello from DataNow!'
```

WHERE THE MINUTES WENT

- 1

Install the SDK

pip install, Python 3.10+

~1 min
- 2

Create a service-account token

Admin-only step in the web client

~5 min
- 3

Set environment variables

Token and workspace name

~1 min
- 4

Run the quickstart

Link a workspace, write, read

~2 min

Time to first call ~10 min

The biggest win is step 2. A developer who is not an admin has to ask someone before writing a line of code. A self-serve trial token would halve the total.

Start with data and ML engineers. The other three audiences wait for their gate.

FIRST AUDIENCE

Data and ML engineers

They already use fsspec-based tools and already search for this problem. Adoption costs one line of code.

- Existing habit: pandas, Dask, PyArrow
- Clear pain: copy-before-compute on large datasets
- Fast proof: a working read in under 5 minutes
- Pulls expansion: training data grows storage inside the account

PHASE 2 • MONTHS 4-6

Platform and DevOps teams

Strong fit for CI/CD artifacts. They expect Terraform or other infrastructure-as-code support first.

REACHED TODAY THROUGH SALES

IT and systems admins

The Management API helps them now. Their buyer already talks to the account team, so the current motion covers them.

PHASE 3 • MONTHS 6-12

Product teams and ISV partners

Partners build once they see adoption. Recruit one or two design partners early so launch day has a real integration.

The strongest alternative is a script that copies data first. DataNow replaces the copy.

	DataNow	Copy-then-process scripts	Cloud S3 mount tools	Review platforms	File transfer services
Many people write to the same files live	● Yes	○ No	○ No edits to existing files	○ Review layer	○ Transfer only
Stream in place, no copy	● Yes	○ Copies first	● Yes	● Finished assets	○ Moves data
File locking	● Yes	○ No	○ No	○ No	○ No
Zero-knowledge encryption built in	● Yes	● Up to you	○ Not built in	○ No	○ No
Programmable from Python	● SDK (beta)	● boto3	● Mount only	● Media API	● API

Four tenets settle arguments before they start.

01

Code comes first

Every page opens with working code. Claims come after the code runs.

02

Fix the path, then drive traffic

No launch push until a new developer reaches a first call in under 5 minutes.

03

One audience at a time

Data and ML engineers first. Each new audience earns its turn by clearing a gate.

04


Sales stays whole

The developer motion feeds enterprise deals. Account teams see every qualified account.

Each use case is a tutorial: the problem, working code, and the result.

1 · Problem The question a developer types into search or an AI assistant	2 · Code Copy, paste, and run in under 5 minutes	3 · Result The number that changed	4 · Next A link to the next tutorial
P1 · FIRST Stream training data to GPU nodes without copying Read shards in place with pandas or PyTorch data loaders. ML engineers	P2 Read and write Parquet with pandas and Dask The one-line path change, shown end to end on a real dataset. Data engineers	P3 Zero-copy ETL from S3 with Connect Link an existing bucket, then process it without a migration. Data and platform	
P4 Publish CI/CD build artifacts from a container Write build outputs straight to a shared workspace. Platform and DevOps	P5 Generate file previews from a headless worker Automate a media pipeline step with no desktop client. File pipelines	P6 Provision workspaces when a project starts Management API scripts for users, groups, and permissions. IT admins	

Three stages, each defined by an event we can log.



`event: first docs or tutorial view`

Discover

Portal, tutorials, PyPI and GitHub listings, answer-engine-ready pages.

`event: first successful SDK file operation`

Activate

Self-serve token and a 5-minute quickstart.

`event: SDK activity in 3 of 4 weeks`

Adopt

Tutorials 2 to 6, office hours, the first case study.

`handoff: product-qualified account`

To the account team

An existing customer with 3+ service accounts or steady SDK traffic for 14 days triggers an alert in the CRM.

By Day 90, a new developer reaches a first call in under 5 minutes, and one customer story is live.

DAY 90

MONTH 1 · LISTEN

MONTH 2 · BUILD

MONTH 3 · SHIP + MEASURE

PMM

owner: me

Friction audit + baselines
done = numbers for every funnel stage

Tutorials 1-3, battlecard, launch plan
done = live and linked from the portal

Docs + DevRel

partner

Code-first quickstart
done = new developer finishes in under 5 min

Product

partner

Self-serve trial token
done = no admin needed for a first call

Sales + CS

partner

Find a design partner
done = signed to a case study

Case study published
done = named customer, measured result

A free way in, and one Tier 1 moment when the SDK reaches GA.

PACKAGING • A RECOMMENDATION

Developer trial workspace

WHAT IT IS

A small free quota with a service account included. No sales call.

VALUE METRIC

Data streamed through the SDK.

PAID PATH

Existing Business and Enterprise plans, which already include service accounts.

VALIDATE WITH

Product and Finance, before the Tier 1 launch.

T1

SDK general availability

Launch post built around code. Portal update. Tutorials 1-3. A design-partner quote. Talk submissions to PyCon and one industry conference.

T2

Each new tutorial or Connect storage type

Changelog, newsletter, social, and a post in the community channels developers already use.

T3

SDK point releases

Changelog and GitHub release notes with migration steps.

Four partner teams, one input each, one monthly review.



Monthly platform business review

Who: PMM, Product, the Sales lead, the CS lead, Partnerships.

Inputs: funnel dashboard, friction log, product-qualified account list.

Decisions: the next tutorial, the next fix, and whether a phase gate is met.

Why expansion credit matters: if reps get no credit when a developer grows an account, they won't work the lead.

Precedent: MongoDB has 53,100+ Atlas customers, most of them self-serve. Customers sold through direct sales drive 88% of subscription revenue.

One north star. Inputs we control, outputs we report.

North star

Weekly active accounts running SDK file operations

Baseline in week 1. Every input below should move it.

INPUTS • WE CONTROL

Time to first call

~10 → <5 min

My hand-timed test is the baseline.

Tutorials published

0 → 5

By Day 90.

Self-serve activation

Instrument

Signup to first file operation, by Day 30.

Docs-to-install rate

Baseline

Set in month 1.

OUTPUTS • WE REPORT

PyPI installs per month

500

Hypothesis for month 6. Reset after baseline.

Case studies

1, then 3

By Day 90, then by month 12.

Product-qualified accounts

Baseline

Routed to account teams each month.

NDR, platform vs. not

+10 pts

Hypothesis for month 12. Control for account size.

What could go wrong, the early signal, and the fix.

RISK	LIKELIHOOD	EARLY SIGNAL	MITIGATION	OWNER
Sales does not follow up on developer-sourced accounts	HIGH	Product-qualified accounts with no follow-up in 7 days	Expansion credit, CRM alerts, review each month	Sales lead
The token step stays admin-only	HIGH	Time to first call stays above 5 minutes	Trial workspace with a token issued at signup	Product
Breaking changes during the SDK beta hurt trust	MED	GitHub issues spike after a release	Version-pinned tutorials, migration notes in every changelog	Product + PMM
Four entry points confuse developers	MED	Tickets that ask which tool to use	One “choose your tool” page and one quickstart per job	Docs
A cloud provider improves its S3 mount tool, or a review platform adds shared storage	LOW	A competitor launch	Lead with multi-writer and locking. Refresh the battlecard within a week.	PMM


TEST BY DAY 30

Developers can start a trial without sales

Python covers the beachhead

SDK GA lands inside the 90 days

Each phase opens only when the last one clears its gate.



One audience first. A fast first call. One customer who proves it. Then expand.

01

Findable

Portal, tutorials, and listings where developers already look.

02

Fast

A first call in under 5 minutes, with no admin in the loop.

03

Proven

A named customer with a measured result by Day 90.

04

Measured

One north star, logged events, and gates that decide what ships.

SOURCES AND NOTES

Where the numbers come from.

Fictional company. DataNow, its products, and its installed-base figures are invented for this sample. The plan is modeled on a real developer-platform launch.

Developer information sources. SlashData, “Where do developers go to stay up-to-date?” Developer Nation survey, Q3 2024, 9,000+ developers.

Documentation use. Stack Overflow Developer Survey 2024, insights for B2B tech marketers.

Time to first call. Postman, “The Most Important API Metric Is Time to First Call,” 2021.

fsspec. The filesystem interface used by pandas, Dask, and PyArrow (filesystem-spec.readthedocs.io).

Overlapping SDKs. Dropbox, “Deprecating the Sync and Datastore APIs,” April 2015.

Self-serve land, sales expand. MongoDB Q4 FY25 results: 53,100+ Atlas customers; direct-sales customers drove 88% of subscription revenue.

First-hand test. My own install and first call on a comparable public Python SDK, June 2026. Timings are hand-measured.

DataNow is a fictional company. Independent work, written in June 2026.